A Journal on Training an LLM from Scratch

Leonard Cseres

October 07, 2025

Contents

Introduction

I starting working on this project with the interest to reproduce what code completion large language models are capable of doing. My goal is to optimize a tiny completions model, and discover where are the limits.

This journal serves as documentation for the design decisions, implementation choices, parameter changes, etc.

Disclaimer: This is an education project and there will be mistakes. Contributions are welcome!

Week 22.09.25

Since we are building a new model from scratch, I wanted to also retrain a new tokenizer. In addition to learning how all this works, I wanted to make changes in the regex parsing, which groups tokens together.

Building the Tokenizer

I decided to reimplement BPE (byte-pair-encoding) tokenization. The main purpose was to learn the BPE algorithm and control it.

I know I needed a high-performance language for this take, so I started out implementing a version in C. However, I quickly realized that the algorithm depends on quite a lot of data structures (vectors, hash maps, heaps) in addition to the regex parsing. Therefore, I switched to C-style C++.

For the training data, I trained the tokenizer exclusively on Python code using the `py150` dataset.

Since python code is *usually* written using ASCII characters, I removed UTF-8 characters before tokenizing, effectively excluding them from the vocabulary. This removes tokens that are very rarely used and allows the model to focus on the essential.

Creating the Regex

I took inspiration from the `cl100k_base` regex [1]:

```
'(?i:[sdmt]|ll|ve|re)|[^\r\n\p{L}\p{N}]?+\p{L}++|\p{N}{1,3}+| ?[^\s\p{L}\p{N}]+
+[\r\n]*+|\s++$|\s*[\r\n]|\s+(?!\S)|\s
```

and ended up on the following regex:

```
 ?[A-Za-z_(][A-Za-z_.]*|%(?:\.\d+)?[sdifFeEgGxXoc%]|[0-9]{1,3}| ?[^ %_A-Za-
z0-9]+(?: ")?[\r\n]*|%|\s+$|\s+(?=\s)|\s
```

- I removed UTF-8 handling and compound expression grouping (ex: `'ve`)
- `?[A-Za-z_(][A-Za-z_.]*` groups together characters with `_` , `(` , and `.` symbols
- `%(?:\.\d+)?[sdifFeEgGxXoc%]` groups together printf formats like `%s` and the rest says very similar

4

Week 29.09.25

Building the Training Loop

I used the `torch` along with `lightning` library [2] to implement a training loop.

For the model, instead of reimplementing from scratch I used the Qwen3 (without MoE) model from the `transformers` library [3].

I choose AdamW for the optimizer with a warm-up and cosine schedule. Below is a non exhaustive list of the parameters chosen:

```
data:
  split_ratio: 0.7
  vocab_size: 20260 # 256 byte tokens + 20000 BPE merges + 4 special tokens
(BOS, EOS, PAD, UNK)

model:
  # Qwen3 architecture parameters
  hidden_size: 512
  num_hidden_layers: 4
  num_attention_heads: 16
  num_key_value_heads: 8
  intermediate_size: 1024
  max_position_embeddings: 2048
  rope_theta: 10000.0
  attention_dropout: 0.1
  rms_norm_eps: 0.000001

training:
  batch_size: 32
  epochs: 50
  lr: 0.0001
  weight_decay: 0.01
  grad_clip: 1.0
  gradient_accumulation_steps: 4
```

Position IDs & Attention Mask

TODO

Week 06.10.25

BOD Special Token

I added this "beginning of sequence" special token to allow it to act as a "attention skin" [4].

The `BOS` token is added at the start of every input sequence to the model. Is is not like the `EOS` (end of sequence) token, where it is added between documents to delimit them.

```
input = [BOS, 234, 6236, 346, 4357, 347, ...]  # where BOS is the token id for
BOS
```

Training Fixes

I updated the scheduler to step on every step instead of epoch. I also scaled the training loss over the accumulated batches instead of using the loss of the step.

Tokenizer Regex

I realized that it cannot hurt to join compound expression grouping. So I added the following capture group:

```
'(?i:[sdmt]|ll|ve|re)
```

Data Distribution

I also removed `max_batches_per_epoch` , replacing it with `max_tokens` such that we have a fair distribute. Before the training batch was shuffled and all data was included but limited to 750. However the evaluation batch was not shuffled but also limited to 750. This meant that the model was training on more that 70% training data, and the validation data was not representative.

PAD Special Token

I've encountered some issues with incomplete batches. Either I drop the last batch or I pad it. I've decided to pad it with the PAD special token. Currently, the model does not see it very often when training so it might not work too well.

Run #1

`commit 2361ffbae2b892f689e9af7984e596809f5462fe`

During this first run the model overfitted as the loss plot shows on Figure 1. This is due to the $10'000'000$ token limit I set, which is too low.


Figure 1: Loss for run #1

Validation Data

One important insight is large language models, when trained on vasts amount of data, have a hard time overfitting.

Therefore, usually we have a much smaller validation set, and instead rely on training metrics. I initially set it so 30% validation, and by increasing the number of tokens, I will decrease the size of the validation set.

Another thing to take into account, since I am randomly splitting data, there might be some similar code ending up in the validation set, contaminating it. So I will have to do something about this.

I also added the perplexity score to track the progress of the model.

Bibliography

[1] OpenAI, "cl100k_base tokenizer regular expression in tiktoken_ext/ openai_public.py, line 89." 2025.

[2] Lightning, "Lightning GitHub." 2025.

[3] Huggingface, "Transformers GitHub." 2025.

[4] G. Xiao, Y. Tian, B. Chen, S. Han, and M. Lewis, "Efficient Streaming Language Models with Attention Sinks." [Online]. Available: https://arxiv. org/abs/2309.17453